




XRootD Client Configuration & API Reference

March 8, 2019

Release 4.9.0 and above

Michal Simon (CERN)



Contents

1	Introduction	3
2	Configuration File	3
3	Command line tools (xrddcp/xrdfs)	4
3.1	xrddcp - copy files	4
3.2	xrdfs - xrootd file and directory meta-data utility	7
3.3	Return Codes	10
4	Environment Variables	10
4.1	Categories	11
4.2	Index of Environment Variables	13
4.2.1	XRD_APPNAME	13
4.2.2	XRD_CLIENTMONITOR	13
4.2.3	XRD_CLIENTMONITORPARAM	13
4.2.4	XRD_CONNECTIONWINDOW	13
4.2.5	XRD_CONNECTIONRETRY	13
4.2.6	XRD_CPCHUNKSIZE	13
4.2.7	XRD_CPINITTIMEOUT	13
4.2.8	XRD_CPPPARALLELCHUNKS	13
4.2.9	XRD_CPTPCTIMEOUT	14
4.2.10	XRD_DATASERVERTTL	14
4.2.11	XRD_GLFNREDIRECTOR	14
4.2.12	XRD_LOADBALANCERTTL	14
4.2.13	XRD_LOCALMETALINKFILE	14
4.2.14	XRD_LOGFILE	14
4.2.15	XRD_LOGLEVEL	14
4.2.16	XRD_LOGMASK	14
4.2.17	XRD_MAXMETALINKWAIT	15
4.2.18	XRD_METALINKPROCESSING	15
4.2.19	XRD_NETWORKSTACK	15
4.2.20	XRD_NODELAY	15
4.2.21	XRD_OPENRECOVERY	16
4.2.22	XRD_PARALLELEVTLOOP	16
4.2.23	XRD_PLUGIN	16
4.2.24	XRD_PLUGINCONFDIR	16
4.2.25	XRD_POLLERPREFERENCE	16
4.2.26	XRD_PREFERIPV4	16
4.2.27	XRD_READRECOVERY	16
4.2.28	XRD_REDIRECTLIMIT	16
4.2.29	XRD_REQUESTTIMEOUT	16
4.2.30	XRD_RUNFORKHANDLER	17
4.2.31	XRD_STREAMERRORWINDOW	17
4.2.32	XRD_STREAMTIMEOUT	17

4.2.33	XRD_SUBSTREAMSPERCHANNEL	17
4.2.34	XRD_TCPKEEPALIVE	17
4.2.35	XRD_TCPKEEPALIVEINTERVAL	17
4.2.36	XRD_TCPKEEPALIVEPROBES	17
4.2.37	XRD_TCPKEEPALIVETIME	17
4.2.38	XRD_TIMEOUTRESOLUTION	18
4.2.39	XRD_WORKERTHREADS	18
4.2.40	XRD_WRITERECOVERY	18
4.2.41	XRD_XCPBLOCKSIZE	18
4.3	Timeouts Explained	18
4.3.1	Connection Window and Connection Retry	18
4.3.2	Stream Timeout	19
4.3.3	Stream Error Window	19
4.3.4	RequestTimeout	19
4.3.5	Time To Live	20
4.3.6	How does it all come together?	20
4.3.7	xrdcp / XrdCl::CopyProcess Third-Party-Copy timeouts	21
5	Client Declarative API	22
5.1	Operation Utilities	22
5.2	Operation Handlers	23
5.3	Pipelining Semantics	25
5.4	List of Operations	26
5.4.1	File Operations	26
5.4.2	FileSystem Operations	30

1 Introduction

This document describes the client (*XrdCl*) component of XRootD framework. In particular it focuses on configuration options of the client (configuration file, environment variables, parameters) and on how they interact with each other.

XrdCl is a multi-threaded C++ implementation of XRootD client based on an *event-loop*, and is provided by the *libXrdCl* library. The standard C++ API is documented [here](#) and is out of the scope of this document. The *XrdCl* implementation is fully asynchronous (all the synchronous calls have been implemented in terms of their asynchronous counterparts). All issued *requests* are queued and sequentially processed by a single-threaded *socket event-loop* (however in order to increase performance it is possible to employ more than one *event-loop*). Also, all incoming *responses* are processed by the *event-loop*, however all the response handlers are executed in a *thread-pool*. The behavior of *XrdCl* can be tuned using a configuration file, environment variables or *XrdCl::DefaultEnv* utility.

Low level connection handling is hidden from the user. Once a request is issued a connection between the client and the server will be established automatically, the connection will be kept alive for further reuse until TTL timeout elapses. By default *XrdCl* is multiplexing all request through a single physical connection, however it is possible to force the component to use multiple physical connections (up to 16) in order to increase the performance over WAN networks. It is also possible to force *XrdCl* to disconnect from a server (e.g. in order to reestablish the connection with a new credential). When the connection between client and server is being established the server may request the client to authenticate (if so, the server will send a list of acceptable authentication methods, e.g. *krb5*, *gsi*, etc.).

The **XrdCl** library is the base for following components: the command line interface (*xrdcp* and *xrdfs*), [python bindings](#), [SSI client](#) and the [Posix API](#).

In addition, this document, in the last section, describes the new declarative client API introduced in **version 4.9.0**.

2 Configuration File

This section describes the XRootD client configuration file. By default XRootD client will use the global config file: */etc/xrootd/client.conf*. However, those settings might be overwritten by the user specific config file: *~/.xrootd/client.conf* and Environment Variables. For the complete list of configurable parameters please consult the [Index of Environment Variables](#).

XRootD client supports protocol- and endpoint-level plug-ins. By convention a single config file is expected per plug-in, as they are discovered and configured by scanning configuration files. The plug-in manager will search for configuration files in:

- */etc/xrootd/client.plugins.d/*,

- `~/.xrootd/client.plugins.d/`,
- and at a location pointed to by: `XRD_PLUGINCONFDIR`

An XRootD client plug-in configuration file should contain following key-value pairs:

- **url** followed by list of endpoints (by default root protocol is assumed) or protocols
- **lib** followed by a path to the library implementing given plug-in
- **enable** followed by *true* or *false*

For example the following config file defines a plug-in for *host.domain.edu* endpoint (root protocol is being assumed) and *http* protocol:

```

1 url = host.domain.edu;http://*
2 lib = /usr/lib64/libAwesomePlugIn.so
3 enabled = true
4
```

3 Command line tools (xrscp/xrdfs)

3.1 xrscp - copy files

xrscp [options] source destination

DESCRIPTION

The **xrscp** utility copies one or more files from one location to another. The data source and destination may be a local or remote file or directory. Additionally, the data source may also reside on multiple servers.

OPTIONS

-C | **--cksum** type [:value | print | source]

Obtains the checksum of type (i.e. `adler32`, `crc32`, or `md5`) from the source, computes the checksum at the destination, and verifies that they are the same. If a value is specified, it is used as the source checksum. When print is specified, the checksum at the destination is printed but is not verified.

-d | **--debug** lvl

Debug level: 1 (low), 2 (medium), 3 (high).

-F | **--coerce**

Ignores locking semantics on the destination file. This option may lead to file corruption if not properly used.

-f | --force

Re-creates a file if it is already present.

-h | --help

Displays usage information.

-H | --license

Displays license terms and conditions.

-N | --nopbar

Does not display the progress bar.

-P | --posc

Requests POSC (persist-on-successful-close) processing to create a new file. Files are automatically deleted should they not be successfully closed.

-D | --proxy proxyaddr:proxyport [NOT YET IMPLEMENTED]

Use proxyaddr:proxyport as a SOCKS4 proxy. Only numerical addresses are supported.

-r | --recursive

Recursively copy all files starting at the given source directory.

--server

Runs as if in a server environment. Used only for server-side third party copy support.

-s | --silent

Neither produces summary information nor displays the progress bar.

-y | --sources num

Uses up to num sources to copy the file.

-S | --streams num

Uses num additional parallel streams to do the transfer. The maximum value is 15. The default is 0 (i.e., use only the main stream).

--tpc [delegate] first | only

Copies the file from remote server to remote server using third-party-copy protocol (i.e., data flows from server to server). The source and destination servers must support third party copies. Additional security restrictions may apply and may cause the copy to fail if they cannot be satisfied. Argument '**first**' tries tpc and if it fails, does a normal copy; while '**only**' fails the copy unless tpc succeeds. When '**delegate**' is specified, the copy delegates the command issuer's credentials to the target server which uses those credentials to authenticate with the source server. Delegation is ignored if the target server is not configured to

use delegated credentials. Currently, only gsi credentials can be delegated.

-v | --verbose

Displays summary output.

-V | --version

Displays version information and immediately exits.

-z | --zip file

Copy given file from a ZIP archive (same as xrdcl.unzip opaque info).

-X | --xrate rate [NOT YET IMPLEMENTED]

Limits the copy speed to the specified rate. The rate may be qualified with the letter **k**, **m**, or **g** to indicate kilo, mega, or giga bytes, respectively. The option only applies when the source or destination is local.

-Z | --dynamic-src

File size may change during the copy.

-I | --infiles fn

Specifies the file that contains a list of input files.

-p | --path

Automatically create remote destination path.

--parallel n

Number of copy jobs to be run simultaneously.

--allow-http

Allow HTTP as source or destination protocol. Requires the XrdClHttp client plugin.

LEGACY OPTIONS

Legacy options are provided for backward compatability. These are now deprecated and should be avoided.

-adler

Equivalent to "**-cksum adler32:source**".

-DI pname numberval

Set the internal parameter pname with the numeric value numberval.

-DS pname stringval

Set the internal parameter pname with the string value stringval.

-md5

Equivalent to "**-cksum md5:source**".

-np

Equivalent to "**-nopbar**".

-OD cgi

Add cgi information cgi to any destination xrootd URL. You should specify the opaque information directly on the destination URL.

-OS cgi

Add cgi information cgi to any source xrootd URL.

-x

Equivalent to "**-sources 12**".

OPERANDS

source: a dash (i.e. -) indicating stanard in, a local file, a local directory name suffixed by /, or an xrootd URL in the form of:

xroot:// [user@] host [:port] /absolutepath

The absolutepath can be a directory.

destination: a dash (i.e. -) indicating stanard out, a local file, a local directory name suffixed by /, or an xrootd URL in the form:

xroot:// [user@] host [:port] /absolutepath

The absolutepath can be a directory.

3.2 xrd fs - xrootd file and directory meta-data utility

xrd fs [-no-cwd] host[:port] [command [args]]

DESCRIPTION

The **xrd fs** utility executes meta-data oriented operations (e.g., ls, mv, rm, etc.) on one or more xrootd servers. Command help is available by invoking xrd fs with no command line options or parameters and then typing "help" in response to the input prompt.

OPTIONS**-no-cwd**

No CWD is being preset in interactive mode.

COMMANDS

chmod path <user><group><other>

Modify permissions of the path. Permission string example: rwxr-x-x

ls [-l] [-u] [-R] [dirname]

Get directory listing.

- l stat every entry and print long listing
- u print paths as URLs
- R list subdirectories recursively
- D show duplicate entries

locate [-n] [-r] [-d] <path>

Get the locations of the path.

- r refresh, don't use cached locations
- n make the server return the response immediately (it may be incomplete)
- d do a recursive, deep locate in order to find data servers
- m prefer host names to IP addresses
- i ignore network dependencies (IPv6/IPv4)

mkdir [-p] [-m<user><group><other>] <dirname>

Creates a directory/tree of directories.

- p create the entire directory tree recursively
- m<user><group><other> permissions for newly created directories

mv <path1> <path2>

Move path1 to path2 locally on the same server.

stat <path>

Get info about the file or directory.

- q query optional flag query parameter that makes xrdfs return error code to the shell if the requested flag combination is not present; flags may be combined together using '|' or '&' Available flags: **XBitSet**, **IsDir**, **Other**, **Offline**, **POSCPending**, **IsReadable**, **IsWriteable**

statvfs <path>

Get info about a virtual file system.

query <code> <params>

Obtain server information. Query codes:

config <**what**> Server configuration; <what> is one of the following:

- bind_max - the maximum number of parallel streams
- checksum - the supported checksum
- pio_max - maximum number of parallel I/O requests

- readv_ior_max - maximum size of a readv element
- readv_iov_max - maximum number of readv entries
- tpc - support for third party copies
- wan_port - the port to use for wan copies
- wan_window - the wan_port window size
- window - the tcp window size
- cms - the status of the cmsd
- role - the role in a cluster
- sitename - the site name
- version - the version of the server

checksumcancel <path> File checksum cancelation

checksum <path> File checksum

opaque <arg> Implementation dependent

opaquefile <arg> Implementation dependent

space <space> Logical space stats

stats <what> Server stats; <what> is a list of letters indicating information to be returned:

- a - all statistics
- p - protocol statistics
- b - buffer usage statistics
- s - scheduling statistics
- d - device polling statistics
- u - usage statistics
- i - server identification
- z - synchronized statistics
- l - connection statistics

xattr <path> Extended attributes

rm <filename>

Remove a file.

rmdir <dirname>

Remove a directory.

truncate <filename> <length>

Truncate a file.

prepare [-c] [-f] [-s] [-w] [-p priority] filenames

Prepare one or more files for access.

- -c co-locate staged files if possible
- -f refresh file access time even if the location is known
- -s stage the files to disk if they are not online
- -w whe files will be accessed for modification
- -p priority of the request, 0 (lowest) - 3 (highest)

cat [-o localfile] file

Print contents of a file to stdout

-o print to the specified local file

tail [-c bytes] [-f] file

Output last part of files to stdout.

-c num.bytes out last num.bytes -f output appended data as file grows

spaceinfo path

Get space statistics for given path.

3.3 Return Codes

- **0** : success
- **50** : generic error (e.g. config, internal, data, OS, command line option)
- **51** : socket related error
- **52** : postmaster related error
- **53** : XRootD related error
- **54** : redirection error
- **55** : query response was negative (this is not an error)

4 Environment Variables

This section describes XRootD client environment variables. The following list of environment variables applies to *xrdcp*, *xrdfs* any other application using the libXrdCl library, unless specified otherwise.

4.1 Categories

Limits/Performance:	<ul style="list-style-type: none">• <code>XRD_PARALLELEVTLOOP</code>• <code>XRD_REDIRECTLIMIT</code>• <code>XRD_SUBSTREAMSPERCHANNEL</code>• <code>XRD_WORKERTHREADS</code>
Logging:	<ul style="list-style-type: none">• <code>XRD_LOGLEVEL</code>• <code>XRD_LOGFILE</code>• <code>XRD_LOGMASK</code>
Metalinks:	<ul style="list-style-type: none">• <code>XRD_METALINKPROCESSING</code>• <code>XRD_LOCALMETALINKFILE</code>• <code>XRD_GLFNREDIRECTOR</code>• <code>XRD_MAXMETALINKWAIT</code>
Monitoring:	<ul style="list-style-type: none">• <code>XRD_APPNAME</code>• <code>XRD_CLIENTMONITOR</code>• <code>XRD_CLIENTMONITORPARAM</code>
Networking:	<ul style="list-style-type: none">• <code>XRD_NETWORKSTACK</code>• <code>XRD_PREFERIPV4</code>
Plug-in:	<ul style="list-style-type: none">• <code>XRD_PLUGIN</code>• <code>XRD_PLUGINCONFDIR</code>

Recovery:	<ul style="list-style-type: none"> • XRD.CONNECTIONRETRY • XRD.OPENRECOVERY • XRD.READRECOVERY • XRD.WRITERECOVERY • XRD.STREAMERRORWINDOW
TCP:	<ul style="list-style-type: none"> • XRD.NODELAY • XRD.TCPKEEPALIVE • XRD.TCPKEEPALIVEINTERVAL • XRD.TCPKEEPALIVEPROBES • XRD.TCPKEEPALIVETIME
Timeouts:	<ul style="list-style-type: none"> • XRD.CONNECTIONWINDOW • XRD.REQUESTTIMEOUT • XRD.STREAMTIMEOUT • XRD.DATASERVERTTL • XRD.LOADBALANCERTTL • XRD.TIMEOUTRESOLUTION
XrCl::CopyProcess / xrdep:	<ul style="list-style-type: none"> • XRD.CPCHUNKSIZE • XRD.CPINITTIMEOUT • XRD.CPTPCTIMEOUT • XRD.CPPARALLELCHUNKS • XRD.XCPBLOCKSIZE
Others:	<ul style="list-style-type: none"> • XRD.POLLERPREFERENCE • XRD.RUNFORKHANDLER

4.2 Index of Environment Variables

4.2.1 XRD_APPNAME

Override the application name reported to the server.

Default: disabled

4.2.2 XRD_CLIENTMONITOR

Path to the client monitor library.

Default: disabled

4.2.3 XRD_CLIENTMONITORPARAM

Additional optional parameters that will be passed to the monitoring object on initialization.

Default: disabled

4.2.4 XRD_CONNECTIONWINDOW

A time window for the connection establishment. A connection failure is declared if the connection is not established within the time window. If a connection failure happens earlier then another connection attempt will only be made at the beginning of the next window.

Default: 120 (seconds)

4.2.5 XRD_CONNECTIONRETRY

Number of connection attempts that should be made (number of available connection windows) before declaring a permanent failure.

Default: 5

4.2.6 XRD_CPCHUNKSIZE

Size of a single data chunk handled by xrdcp / XrdCl::CopyProcess.

Default: 16KiB

4.2.7 XRD_CPINITTIMEOUT

Maximum time allowed for the copy process to initialize, ie. open the source and destination files.

Default: 600 (seconds)

4.2.8 XRD_CPPARALLELCHUNKS

Maximum number of asynchronous requests being processed by the xrdcp / XrdCl::CopyProcess command at any given time.

Default: 4

4.2.9 XRD_CPTPCTIMEOUT

Maximum time allowed for a third-party copy operation to finish.

Default: 1800 (seconds)

4.2.10 XRD_DATASERVERTTL

Time period after which an idle connection to a data server should be closed.

Default: 300 (seconds)

4.2.11 XRD_GLFNREDIRECTOR

The redirector will be used as a last resort if the GLFN tag is specified in a Metalink file.

Default: none

4.2.12 XRD_LOADBALANCERTTL

Time period after which an idle connection to a manager or a load balancer should be closed.

Default: 1200 (seconds)

4.2.13 XRD_LOCALMETALINKFILE

Enable/Disable local Metalink file processing (by convention the following URL schema has to be used: *root://localfile//path/filename.meta4*) The *localfile* semantic is now deprecated, use *file://localhost/path/filename.meta4* instead!

Default: 0

4.2.14 XRD_LOGFILE

If set, the diagnostics will be printed to the specified file instead of *stderr*.

Default: disabled

4.2.15 XRD_LOGLEVEL

Determines the amount of diagnostics that should be printed. Valid values are: Dump, Debug, Info, Warning, and Error.

Default: disabled

4.2.16 XRD_LOGMASK

Determines which diagnostics topics should be printed at all levels. It's a "|" separated list of topics. The first element may be "All" in which case all the topics are enabled and the subsequent elements may turn them off, or "None" in which case all the topics are disabled and the subsequent flags may turn them on. If the topic name is prefixed with "^", then it means that the topic should

be disabled. If the topic name is not prefixed, then it means that the topic should be enabled.

The log mask may as well be handled for each diagnostic level separately by setting one or more of the following variables: XRD_LOGMASK_ERROR, XRD_LOGMASK_WARNING, XRD_LOGMASK_INFO, XRD_LOGMASK_DEBUG, and XRD_LOGMASK_DUMP.

Available topics: AppMsg, UtilityMsg, FileMsg, PollerMsg, PostMasterMsg, XRootDTransportMsg, TaskMgrMsg, XRootDMsg, FileSystemMsg, AsyncSockMsg

Default: The default for each level is "All", except for the Dump level, where the default is "All ^PollerMsg". This means that, at the Dump level, all the topics but "PollerMsg" are enabled.

4.2.17 XRD_MAXMETALINKWAIT

The maximum time in seconds a client can be stalled by the server if a Metalink redirector is available.

Default: 60 (seconds)

4.2.18 XRD_METALINKPROCESSING

Enable/Disable Metalink processing.

Default: 1

4.2.19 XRD_NETWORKSTACK

The network stack that the client should use to connect to the server. Possible values are:

- **IPAuto** - automatically detect which IP stack to use
- **IPAll** - use IPv6 stack (AF_INET6 sockets) and both IPv6 and IPv4 (mapped to IPv6) addresses
- **IPv6** - use only IPv6 stack and addresses
- **IPv4** - use only IPv4 stack (AF_INET sockets) and addresses
- **IPv4Mapped6** - use IPv6 stack and mapped IPv4 addresses

Default: IPAuto

4.2.20 XRD_NODELAY

Disables the Nagle algorithm if set to 1 (default), enables it if set to 0.

Default: 1

4.2.21 XRD_OPENRECOVERY

Determines if open recovery should be enabled or disabled for mutable (truncate or create) opens.

Default: `true`

4.2.22 XRD_PARALLELEVTLOOP

The number of event loops.

Default: `1`

4.2.23 XRD_PLUGIN

A default client plug-in to be used.

Default: `none`

4.2.24 XRD_PLUGINCONFDIR

A custom location containing client plug-in config files.

Default: `none`

4.2.25 XRD_POLLERPREFERENCE

A comma separated list of poller implementations in order of preference.

Default: `built-in`

4.2.26 XRD_PREFERIPV4

If set the client tries first IPv4 address (turned off by default).

Default: `0`

4.2.27 XRD_READRECOVERY

Determines if read recovery should be enabled or disabled.

Default: `true`

4.2.28 XRD_REDIRECTLIMIT

Maximum number of allowed redirections.

Default: `16`

4.2.29 XRD_REQUESTTIMEOUT

Default value for the time after which an error is declared if it was impossible to get a response to a request.

Default: `1800` (seconds)

4.2.30 XRD_RUNFORKHANDLER

Determines whether the fork handlers should be enabled, making the API fork safe.

Default: 0

4.2.31 XRD_STREAMERRORWINDOW

Time after which the permanent failure flags are cleared out and a new connection may be attempted if needed.

Default: 1800

4.2.32 XRD_STREAMTIMEOUT

Default value for the time after which a connection error is declared (and a recovery attempted) if there are unfulfilled requests and there is no socket activity or a registered wait timeout.

Default: 60 (seconds)

4.2.33 XRD_SUBSTREAMSPERCHANNEL

Number of streams per session.

Default: 1

4.2.34 XRD_TCPKEEPALIVE

Enable/Disable the TCP keep alive functionality.

Default: 0

4.2.35 XRD_TCPKEEPALIVEINTERVAL

Interval between subsequent keepalive probes (Linux only).

Default: 75

4.2.36 XRD_TCPKEEPALIVEPROBES

Number of unacknowledged probes before considering the connection dead (Linux only).

Default: 9

4.2.37 XRD_TCPKEEPALIVETIME

Time between last data packet sent and the first keepalive probe (Linux only).

Default: 7200

4.2.38 XRD_TIMEOUTRESOLUTION

Resolution for the timeout events. Ie. timeout events will be processed only every XRD_TIMEOUTRESOLUTION seconds.

Default: 15 (seconds)

4.2.39 XRD_WORKERTHREADS

Number of threads processing user callbacks.

Default: 3

4.2.40 XRD_WRITERECOVERY

Determines if write recovery should be enabled or disabled.

Default: true

4.2.41 XRD_XCPBLOCKSIZE

Maximu size of a data block assigned to a single source in case of an extreme copy transfer.

Default: 128MiB

4.3 Timeouts Explained

4.3.1 Connection Window and Connection Retry

The *ConnectionWindow* parameter is applied during client-server connection and controls two aspects of this process:

- First of all, it controls the length of time allowed to establish an XRootD connection (physical connection, XRootD hand-shake, and authentication if required). It is important to note that ***Connection Window is applied per physical address***. This means that if a connection fails before the end of current *ConnectionWindow* and another physical address is available it will be tried immediately.
- Secondly (if there are no more available physical addresses), it defines the length of time that must elapse after a connection failure before the connection can be retried. More precisely, the client has to wait until the end of the current *ConnectionWindow* before attempting another connection. The number of retries that might be attempted is controlled by *ConnectionRetry* environment variable.

Let us illustrate all this with following example. Suppose XRootD client wants to connect to a server with three physical IP address (2x IPv6 and 1x IPv4). For the sake of argument let us suppose it will fail after 60s during the hand-shake procedure, while connecting to the 1st IPv6 address. What will happen next? Since there are two more addresses available, the client will immediately proceed to the next one. Now let us suppose that the cumulative time spent

on establishing the physical connection and on carrying out the hand-shake exceeded 120s (nominal value of *ConnectionWindow*). In this case the second connection attempt will be timed out, and XRootD client will proceed to the 3rd IP address. Again, let us suppose that similarly as in case of the 1st IP address the connection failed after 60s. Since there are no more address to try, the client will have to wait until the end of the current *ConnectionWindow* (that is for another 60s) before the connection procedure can be restarted. Now how all this relates to the *ConnectioRetry*? The nominal value of *ConnectioRetry* is 5, which means we can retry the whole procedure four more times (Note: ***ConnectionRetry* is not applied per single physical connection but rather to the whole connection procedure**).

4.3.2 Stream Timeout

The *StreamTimeout* parameter is applied during every request/response exchange after the client and the server established a connection. It defines the maximum length of time that may elapse between the moment when the client has sent a request and the moment when the client has received a response for the request in question. If the time spent waiting for response from the server exceeds the *StreamTimeout* an error is declared (and the client will disconnect form the server).

There are two exceptions to the above stated rule:

- The server may force the client to reissue the request by sending *kXR-wait* response. In this case *StreamTimeout* does not apply to the original request anymore.
- The server may explicitly instruct the client to not apply the *StreamTimeout* to given request by sending *kXR-waitresp*.

4.3.3 Stream Error Window

The *StreamErrorWindow* controls the length of time that needs to elapse after a fatal error before the client may attempt to reconnect to the server. A fatal error is declared eg. if the host name cannot be resolved, a low level Posix system call fails (eg. *connect/fcntl/epool*), or client runs out of connection retries.

4.3.4 RequestTimeout

The *RequestTimeout* parameter is applied to a logical XRootD operation (eg. opening a file, listing directory, etc.) as a whole. It is the maximum length of time that may elapse from the moment an operation has been issued using XRootD client API until it has been resolved (no matter how many underlying requests it will trigger). If the *RequestTimeout* is exceeded an error is declared and the operation is resolved as failed.

Note: The value of this parameter might be overwritten directly by the user of XRootD client API by setting the timeout argument.

4.3.5 Time To Live

A Time To Live (TTL) timeout controls the lifetime of an idle physical connection. If for the given communication channel the time length elapsed from last exchange of request/response between the client and server exceeds the TTL timeout the given connection will be terminated. There are two types of TTL timeouts in XRootD client:

- *DataServerTTL*: a TTL timeout that is applied to Data Servers
- *LoadBalancerTTL*: a TTL timeout that is applied to Managers


4.3.6 How does it all come together?

Let us now consider an example in order to illustrate how all those timeouts play along (**for clarity please consult the diagram below**). Suppose that an *XrdCl::File::Open(...)* operation is being called and that there is no open connection between the client and the server. The client will have to establish the XRootD connection first (subject to *Connection Window*):

- open physical connection
- carry out hand-shake procedure (*kXR_protocol*, *kXR_login*, etc.)

Subsequently, the client will issue a *kXR_open* request (subject to *StreamTimeout*). Let us suppose that the server will respond with a *kXR_redirect* redirecting the client to a data server. In this case, the client will have to open another XRootD connection (again, subject to *Connection Window*) and then send an open request (again, subject to *StreamTimeout*). Finally, once the server responds, the open operation will be resolved. The whole process described in the scenario above is subject to *RequestTimeout*.

Once the connections to the manager and data server become idle they will be subject to respective *TTL timeouts*.



4.3.7 xrdcp / XrdCl::CopyProcess Third-Party-Copy timeouts

The *CPInitTimeout* parameter is applied during initialization of a Third-Party-Copy (TPC) transfer. It defines the maximum length of time that may elapse until TPC transfer has been initialized (ie. *open* destination, *open* source, issue *sync*, for more details please consult the [TPC Protocol Reference](#)).

The *CPTPCTimeout* parameter defines the maximum length of time that may elapse between the moment when the actual transfer has been started and the moment when it is finished (ie. it is applied to the second *sync*, for more details please consult the [TPC Protocol Reference](#)).

5 Client Declarative API

This section describes XRootD client declarative API introduced in version 4.9.0. For the standard *XrdCl::File* and *XrdCl::FileSystem* API please consult our [Doxygen](#) documentation. Similarly as the standard *XrdCl* API, the declarative API allows to issue *File* and *FileSystem* operations, however its sole focus is on facilitating the asynchronous programming model and chaining of operations. Also, the new API has been designed to be more in line with modern C++ programming practices (see example below).

```
1  File f;
2  std::future<ChunkInfo> resp;
3
4  // open, read from and close the file
5  Pipeline p = Open( file , url , OpenFlags::Read)
6              | Read( file , offset , size , buffer ) >> resp
7              | Close( file );
8
9  auto status = WaitFor(p);
10
11
```

5.1 Operation Utilities

There are several utilities for facilitating composition of operations:

- **Pipeline**: a class that can wrap any kind of operation (including compound operations).
- **Async**: a utility for asynchronous execution of pipelines.

Returns: *std::future<XrdCl::XRootDStatus>*

```
1  FileSystem fs( url );
2  std::future<XRootDStatus> status =
3      Async( Truncate( fs , path , size ) );
4
5
```

- **WaitFor**: a utility for synchronous execution of pipelines.

Returns: *XrdCl::XRootDStatus*.

```
1  FileSystem fs( url );
2  XRootDStatus status = WaitFor( Truncate( fs , path , size ) );
3
4
```

- **XrdCl::Fwd**: a *forward* is used to pass values between different operations in a pipeline. In particular it can be used to forward a value from an operation handler to a subsequent operation as an argument.

Consider following example of reading a whole file of unknown size:

```

1  Fwd<uint32_t> size; // size of the file
2  Fwd<void*> buff; // buffer for the data
3
4
5  std::future<ChunkInfo> resp; // server response
6
7  auto &&p = Open( file , url , OpenFlags::Read ) >>
8      [size , buff]( XRootDStatus &status , StatInfo &info )
9      {
10         if (!status.IsOK()) return;
11         size = info.GetSize(); // forward size and
12         buff = new char[info.GetSize()]; // buffer
13     }
14     | Read( file , 0 , size , buff ) >> resp
15     | Close( file );
16
17  auto status = WaitFor( p );
18

```

In lines 2-3 we declare forwardable *size* and *buffer* arguments. In the pipeline we first issue an open, which we handle with a lambda (open returns also stat information). Inside of the lambda (lines 11-12) we set the values of *size* and *buffer*. In the subsequent *Read* (line 14) operation we use the *size* and *buffer* although they values will be only set once we get the response for the preceding *Open*.

- **XrdCl::Parallel:** aggregates several operations (might be compound operations) for parallel execution; as an argument accepts **variable number of operations** or a **container of operations** (see example below).

```

1
2  auto &&o1 = Open( file1 , url1 , OpenFlags::Read );
3  auto &&o2 = Open( file2 , url2 , OpenFlags::Read );
4  auto &&o2 = Open( file2 , url2 , OpenFlags::Read );
5
6  // open 3 files in parallel
7  Pipeline p = Parallel(o1,o2,o3);
8
9  auto status = WaitFor( p );
10

```

5.2 Operation Handlers

The declarative API supports following handlers: *XrdCl::ResponseHandler*, functions, function objects, lambdas, *std::future* and *std::package_task* (consult the list below for respective examples). Each operation defines its **response type** (see [List of Operations](#)) that should be used when constructing a respective handler for the operation.

- **XrdCl::ResponseHandler** – standard XRootD response handler. Operations can accept *XrdCl::ResponseHandler* both by reference and pointer.


```

1
2  class ExampleHandler : public ResponseHandler
3  {
4      public:
5          void HandleResponse(
6              XRootDStatus *status, // status of the operation
7              AnyObject *response // server response (type erased)
8          )
9          {
10             // handle the operation here
11         }
12     };
13
14     ...
15
16     FileSystem fs(url);
17     ExampleHandler hndl;
18
19     auto status = WaitFor( Stat(fs,path) >> hndl );
20

```

- **functions / function objects / lambdas** – the XRootD declarative API plays well with standard *C++ callable elements*. The callback signature has to match:

- `std::function<void(XRootDStatus&)>` for operations that define their response type as *void*.
- `std::function<void(XRootDStatus&, Response&)>` where *Response* is defined as a response type for given operation.

```

1
2  void ExampleHandler(
3      XRootDStatus &status, // status of the operation
4      StatInfo      &info    // server response (explicit type)
5  )
6  {
7      // handle the operation here
8  }
9
10     ...
11
12     FileSystem fs(url);
13     // could also be a lambda or function object !!!
14     auto status = WaitFor( Stat(fs,path) >> ExampleHandler );
15

```

- **`std::future`** – the *future*'s template parameter has to match the response type of given operation. In case of a failure the *future* will throw an instance of `XrdCl::PipelineException` that in turn will yield the `XrdCl::XRootDStatus`.

```

1  File file;
2  std::future<ChunkInfo> resp;
3
4  Async( Read(off, size, buff) >> resp );
5
6
7
8  ...
9
10 // later on process the future
11 try
12 {
13     ...
14     // if everything went OK we will get the ChunkInfo,
15     // otherwise it will throw
16     ChunkInfo chunk = resp.get();
17     ...
18 }
19 catch( PipelineException &ex )
20 {
21     // we will learn the reason for failing from
22     // the status object
23     XRootDStatus &status = ex.GetError();
24     ...
25 }
26

```

- ***std::packaged_task*** is a combination of a lambda and a *std::future*, e.g. it can be used to parse the response with a lambda into a desired type of *std::future*.

```

1
2  using namespace std;
3
4  packaged_task<uint64_t(XRootDStatus &st, StatInfo &info)> parse =
5      [](XRootDStatus &st, StatInfo &info)
6      {
7          if(!st.IsOK) throw PipelineException(st);
8          return info.GetSize();
9      };
10
11  FileSystem fs(url);
12  future<uint64_t> size = parse.get_future();
13
14  Async( Stat(fs, path) >> parse );
15
16 // later on use size the same way as
17 // the future from previous example
18

```

5.3 Pipelining Semantics

Operations can be pipelined using *operator|*. In order to illustrate the pipelining semantics we will consider following scenario: suppose one wants to read 1KB

form a files, however the prerequisite for reading is creating a lock file. Now let us consider following code:

```

1  File lock , file ;
2  FileSystem fs ( url );
3
4  std :: future < ChunkInfo > resp ; // server response
5
6  auto &&p = Open ( lock , " root : // host // path / to / . lock " , OpenFlags :: New )
7             | Close ( lock )
8             | Open ( file , " root : // host // path / to / file . txt " , OpenFlags :: Read )
9             | Read ( file , 0 , 1024 , buff ) >> resp
10            | Close ( file ) ;
11            | Rm ( fs , " root : // host // path / to / . lock " ) ;
12
13 // we can already pass resp to an algorithm for processing
14
15 // we wait for the pipeline to complete
16 auto status = WaitFor ( p ) ;
17

```

In lines 6-7 the lock file is being created. Afterwards, in lines 8-10 the pipeline continues: it does an open, a read and a close on the actual file. Finally, in line 11 the lock file is being deleted. Note that if an operation on the pipeline fails subsequent operations in the pipeline wont be executed, however their handlers will be called with an error status of *errPipelineFailed* (in order to allow for a clean up if necessary). Using the pipelining API makes the source code more coherent and the control flow more explicit.

5.4 List of Operations

There are two types of operations: the *XrdCl::File* operations and *XrdCl::FileSystem* operations. Each operation has a well defined set of arguments, however **any argument might be lifted** to a *std::future* or a *XrdCl::Fwd*. It is possible (**but not mandatory**) to specify a handler for each operation using the streaming operator (*operator>>*). All Operations are non-copyable objects (*move* only).

5.4.1 File Operations

All arguments of any File Operation (except for the *XrdCl::File* object itself) are **liftable to *XrdCl::Fwd* and *std::future***.

- **Open** – open remote / local file

Signature:

- *Open(XrdCl::File &file, ...);*
- *Open(XrdCl::File *file, ...);*

Arguments (remaining):

- url – base **type**: *std::string*
- flags – base **type**: *XrdCl::OpenFlags::Flags*
- mode – base **type**: *XrdCl::Access::Mode*, **default**: *Access::None*

Operation **status**: *XRootDStatus*

Response:

- *void*
- *XrdCl::StatInfo* (not for XrdCl::ResponseHandler)

- **Read** – read data from remote / local file

Signature:

- *Read(XrdCl::File &file, ...);*
- *Read(XrdCl::File *file, ...);*

Arguments (remaining):

- offset – base **type**: *uint64_t*
- size – base **type**: *uint32_t*
- buffer – base **type**: *void**

Operation **status**: *XRootDStatus*

Response: *XrdCl::ChunkInfo*

- **Close** – close remote / local file

Signature:

- *Close(XrdCl::File &file);*
- *Close(XrdCl::File *file);*

Operation **status**: *XRootDStatus*

Response: *void*

- **Stat** – stat the remote / local file

Signature:

- *Stat(XrdCl::File &file, ...);*
- *Stat(XrdCl::File *file, ...);*

Arguments (remaining):

- force – base **type**: *bool*

Operation **status**: *XRootDStatus*

Response: *XrdCl::StatInfo*

- **Write** – write data to remote / local file

Signature:

- *Write(XrdCl::File &file, ...);*
- *Write(XrdCl::File *file, ...);*

Arguments (remaining):

- offset – base **type**: *uint64_t*
- size – base **type**: *uint32_t*
- buffer – base **type**: *void**

Operation **status**: *XRootDStatus*

Response: *void*

- **Sync** – sync the remote / local file

Signature:

- *Sync(XrdCl::File &file);*
- *Sync(XrdCl::File *file);*

Operation **status**: *XRootDStatus*

Response: *void*

- **Truncate** – truncate the remote / local file

Signature:

- *Truncate(XrdCl::File &file, ...);*
- *Truncate(XrdCl::File *file, ...);*

Arguments (remaining):

- size – base **type**: *uint64_t*

Operation **status**: *XRootDStatus*

Response: *void*

- **VectorRead** – vector-read data from remote / local file

Signature:

- *VectorRead(XrdCl::File &file, ...);*
- *VectorRead(XrdCl::File *file, ...);*

Arguments (remaining):

- chunks – base **type**: *XrdCl::ChunkList*
- buffer – base **type**: *void**

Operation **status**: *XRootDStatus*

Response: *XrdCl::ChunkList*

- **VectorWrite** – vector-write data to remote / local file

Signature:

- *VectorWrite(XrdCl::File &file, ...);*
- *VectorWrite(XrdCl::File *file, ...);*

Arguments (remaining):

- chunks – base **type**: *XrdCl::ChunkList*

Operation **status**: *XRootDStatus*

Response: *void*

- **WriteV** – writev data to remote / local file

Signature:

- *WriteV(XrdCl::File &file, ...);*
- *WriteV(XrdCl::File *file, ...);*

Arguments (remaining):

- offset – base **type**: *uint64_t*
- iov – base **type**: *struct iovec**
- iovcnt – base **type**: *int*

Operation **status**: *XRootDStatus*

Response: *void*

- **Fcntl** – issue fcntl for remote / local file

Signature:

- *Fcntl(XrdCl::File &file, ...);*
- *Fcntl(XrdCl::File *file, ...);*

Arguments (remaining):

- arg – base **type**: *XrdCl::Buffer*

Operation **status**: *XRootDStatus*

Response: *XrdCl::Buffer*

- **Visa** – issue fcntl for remote / local file

Signature:

- *Visa(XrdCl::File &file);*
- *Visa(XrdCl::File *file);*

Operation **status**: *XRootDStatus*

Response: *XrdCl::Buffer*

5.4.2 FileSystem Operations

All arguments of any FileSystem Operation (except for the *XrdCl::FileSystem* object itself) are **liftable to *XrdCl::Fwd* and *std::future***.

- **Locate** – locate remote file

Signature:

- *Locate(XrdCl::FileSystem &file, ...);*
- *Locate(XrdCl::FileSystem *file, ...);*

Arguments (remaining):

- path – base **type**: *std::string*
- flags – base **type**: *XrdCl::OpenFlags::Flags*

Operation **status**: *XRootDStatus*

Response: *LocationInfo*

- **DeepLocate** – recursively locate remote file

Signature:

- *DeepLocate(XrdCl::FileSystem &file, ...);*
- *DeepLocate(XrdCl::FileSystem *file, ...);*

Arguments (remaining):

- path – base **type**: *std::string*
- flags – base **type**: *XrdCl::OpenFlags::Flags*

Operation **status**: *XRootDStatus*

Response: *LocationInfo*

- **Mv** – move remote file

Signature:

- *Mv(XrdCl::FileSystem &file, ...);*
- *Mv(XrdCl::FileSystem *file, ...);*

Arguments (remaining):

- path1 – base **type**: *std::string*
- path2 – base **type**: *std::string*

Operation **status**: *XRootDStatus*

Response: *void*

- **Query** – query remote server

Signature:

- *Query(XrdCl::FileSystem &file, ...);*
- *Query(XrdCl::FileSystem *file, ...);*

Arguments (remaining):

- queryCode – base **type**: *XrdCl::QueryCode::Code*
- argument – base **type**: *XrdCl::Buffer*

Operation **status**: *XRootDStatus*

Response: *XrdCl::Buffer*

- **Truncate** – truncate remote file

Signature:

- *Truncate(XrdCl::FileSystem &file, ...);*
- *Truncate(XrdCl::FileSystem *file, ...);*

Arguments (remaining):

- path – base **type**: *std::string*
- size – base **type**: *uint64_t*

Operation **status**: *XRootDStatus*

Response: *void*

- **Rm** – remove remote file

Signature:

- *Rm(XrdCl::FileSystem &file, ...);*
- *Rm(XrdCl::FileSystem *file, ...);*

Arguments (remaining):

- path – base **type**: *std::string*

Operation **status**: *XRootDStatus*

Response: *void*

- **MkDir** – create remote directory

Signature:

- *MkDir(XrdCl::FileSystem &file, ...);*
- *MkDir(XrdCl::FileSystem *file, ...);*

Arguments (remaining):

- path – base **type**: *std::string*

Operation **status**: *XRootDStatus*

Response: *void*

- **RmDir** – remove remote directory

Signature:

- *RmDir(XrdCl::FileSystem &file, ...);*
- *RmDir(XrdCl::FileSystem *file, ...);*

Arguments (remaining):

- path – base **type**: *std::string*

Operation **status**: *XRootDStatus*

Response: *void*

- **ChMod** – change access mode on a remote directory or file

Signature:

- *ChMod(XrdCl::FileSystem &file, ...);*
- *ChMod(XrdCl::FileSystem *file, ...);*

Arguments (remaining):

- path – base **type**: *std::string*
- mode – base **type**: *XrdCl::Access::Mode*

Operation **status**: *XRootDStatus*

Response: *void*

- **Ping** – ping remote server

Signature:

- *Ping(XrdCl::FileSystem &file);*
- *Ping(XrdCl::FileSystem *file);*

Operation **status**: *XRootDStatus*

Response: *void*

- **Stat** – stat remote directory or file

Signature:

- *Stat(XrdCl::FileSystem &file, ...);*
- *Stat(XrdCl::FileSystem *file, ...);*

Arguments (remaining):

- path – base **type**: *std::string*

Operation **status**: *XRootDStatus*

Response: *XrdCl::StatInfo*

- **StatVFS** – status information for a Virtual File System

Signature:

- *StatVFS(XrdCl::FileSystem &file, ...);*
- *StatVFS(XrdCl::FileSystem *file, ...);*

Arguments (remaining):

- path – base **type**: *std::string*

Operation **status**: *XRootDStatus*

Response: *XrdCl::StatInfoVFS*

- **Protocol** – obtain server protocol information

Signature:

- *Protocol(XrdCl::FileSystem &file);*
- *Protocol(XrdCl::FileSystem *file);*

Operation **status:** *XRootDStatus*

Response: *XrdCl::ProtocolInfo*

- **DirList** – list remote directory

Signature:

- *DirList(XrdCl::FileSystem &file, ...);*
- *DirList(XrdCl::FileSystem *file, ...);*

Arguments (remaining):

- path – base **type:** *std::string*
- flags – base **type:** *XrdCl::DirListFlags::Flags*

Operation **status:** *XRootDStatus*

Response: *XrdCl::DirectoryList*

- **SendInfo** – send info to remote server

Signature:

- *SendInfo(XrdCl::FileSystem &file, ...);*
- *SendInfo(XrdCl::FileSystem *file, ...);*

Arguments (remaining):

- info – base **type:** *std::string*

Operation **status:** *XRootDStatus*

Response: *XrdCl::Buffer*

- **Prepare** – prepare one or more files for access

Signature:

- *Prepare(XrdCl::FileSystem &file, ...);*
- *Prepare(XrdCl::FileSystem *file, ...);*

Arguments (remaining):

- fileList – base **type:** *std::vector<std::string>*
- flags – base **type:** *XrdCl::PrepareFlags::Flags*
- priority – base **type:** *uint8_t*

Operation **status:** *XRootDStatus*

Response: *XrdCl::Buffer*